

[← Back to Home](#)

Project 3: [Auto]Stitching and Photo Mosaics

Part A

A.1 Shoot the Pictures

Grimes 1



Taken at the Grimes Engineering Center to the left (with a lamppost in frame)

Grimes 2



Taken at the Grimes Engineering Center to the right showing more of the driveway

Inside Grimes 1



Taken on an upper floor at Grimes, featuring more of the staircase

Inside Grimes 2



Taken on an upper floor at Grimes, with more of the seating in view

I took these photos with the help of my friend and a tripod so that we could minimize any movement that could potentially change the center of projection (COP) and thereby only rotate the camera!

A.2 Recover Homographies

To set up the system of equations for recovering the unknown homography matrix, I first wrote the given terms we have ($im1_pts$, $im2_pts$ as the x , y coordinates from our 2 images) in the form of $Ah = b$. Here, it was particularly important for me to remember that the vector h only has 8 unknowns and not 9, since there are 8 degrees of freedom in this transformation. From there, if we multiply both the known and unknown terms, and then divide by our third coordinate as seen in the steps that I took in the image below, then we have two equations of x' and y' . However, we can rewrite this further by multiplying both sides of both equations and then expanding the left-hand-side (LHS). Finally, we're left with a system of linear equations to solve for the unknown h vector.

Homography Equations

proj 3: stitching photo mosaics

10/2/25

A.2 recover homographies

$$P' = H P, \quad P = (x, y, 1)^T = \text{coordinate in img1}$$

$$P' = (x', y', 1)^T = \text{coordinate in img2}$$

$$AP = B$$

$H = 3 \times 3$ matrix mapping 1 image plane $\rightarrow$ another plane under perspective projection (coef = 8, -1 because scale doesn't matter)

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} \sim \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \left\{ \begin{array}{l} x' = \frac{h_{11}x + h_{12}y + h_{13}}{h_{31}x + h_{32}y + 1} \\ y' = \frac{h_{21}x + h_{22}y + h_{23}}{h_{31}x + h_{32}y + 1} \end{array} \right. \begin{array}{l} \text{divide} \\ \text{by 3rd} \\ \text{term} \\ \text{(to convert} \\ \text{back to 2D} \\ \text{plane)} \end{array}$$

constant,
only 8 dof

$$x'(h_{31}x + h_{32}y + 1) = h_{11}x + h_{12}y + h_{13}$$

$$y'(h_{31}x + h_{32}y + 1) = h_{21}x + h_{22}y + h_{23}$$

$$x'h_{31}x + x'h_{32}y + x' = h_{11}x + h_{12}y + h_{13}$$

$$y'h_{31}x + y'h_{32}y + y' = h_{21}x + h_{22}y + h_{23}$$

expanded LHS terms

$$\begin{cases} x' = h_{11}x + h_{12}y + h_{13} - x'h_{31}x - x'h_{32}y \\ y' = h_{21}x + h_{22}y + h_{23} - y'h_{31}x - y'h_{32}y \end{cases}$$

unknown h terms on RHS, system of eqs. to find x', y'


$$AH = b$$

$$\begin{bmatrix} x & y & 1 & 0 & 0 & 0 & -xx' & -yy' \end{bmatrix} h = x'$$

$$\begin{bmatrix} 0 & 0 & 0 & x & y & 1 & -xx' & -yy' \end{bmatrix} h =$$

How I derived the system of equations used for recovering our homography matrix

Grimes




Correspondences (5 points) visualized on Grimes images

Grimes Recovered Homography

```
[[ 1.60904060e+00 -2.20491579e-01 -3.01711667e+02]
 [ 3.70956942e-01  1.25978761e+00 -1.22533018e+02]
 [ 1.14276180e-03 -5.01873778e-04  1.00000000e+00]]
```

Inside Grimes



Correspondences (5 points) visualized on Grimes images from indoors

Inside Grimes Recovered Homography

```
[ [ 1.34953238e+00 -1.02735203e-01 -2.84574897e+02]
  [-6.49643399e-02  1.20693002e+00  1.95695271e+01]
  [ 5.01632297e-04 -2.92740559e-04  1.00000000e+00] ]
```

A.3 Warp the Images

Now that we've recovered our homography matrix, we can apply the homography to warp our input images. We tried 2 interpolation methods: nearest neighbor and bilinear interpolation. Whereas nearest neighbor interpolation will round coordinates to the nearest pixel value, bilinear interpolation will use a weighted average from the 4 nearest pixels.


The tips from staff included in our project spec were very helpful! In order to handle predicting the bounding box for the resulting warped image, I implemented a helper function that essentially applied the homography transformation to each of the 4 corners of the input image, and then took the minimum and maximum x and y values to basically serve as a container with dimensions to fit the output

warped image. From there, I called this helper function to compute the output bounds in both nearest neighbor and bilinear interpolation. My process involved looping through each coordinate within the bounds of the output image dimensions (computed by our helper) and then mapping that coordinate back into the source image by using the inverse of the homography matrix.

For nearest neighbor interpolation, the source coordinates were rounded to the nearest integer to find the nearest neighbor pixel value, and then that value was assigned to the output pixel. However, for bilinear interpolation, I used the floor and ceiling method on the source coordinates to find the 4 closest pixels. Then, I subtracted these coordinates from the source coordinates to find the distance between, and use this distance to weight the pixel values and take a weighted average first of the top two pixels, then the bottom two pixels, and then finally between the top and bottom averages to get the final weighted average pixel value to assign to the output pixel.


When it comes to comparing quality and speed, nearest neighbor interpolation ran faster than bilinear interpolation, but as seen in my images below (may need to zoom in a bit!), bilinear interpolation tended to produce a more smooth warping in tradeoff for speed. When testing my code by trying image rectification, nearest neighbor interpolation took around 5.6s whereas bilinear interpolation took around 9s on my image of a poke bowl, and likewise nearest neighbor vs. bilinear interpolation ran for 13.1s vs. 21.6s for my image of bagels.

Warped Poke Bowl



Comparison of nearest neighbor vs. bilinear interpolation on a poke bowl

Warped Bagel



Comparison of nearest neighbor vs. bilinear interpolation on bagels

A.4 Blend the Images into a Mosaic


To create an image mosaic, I tried implementing the one-shot procedure by warping all images at once. However, I realized I had to modify my bilinear interpolation function to handle the entire output shape of the mosaic, rather than

try to combine and blend images without knowing the dimensions of the output. I also tried using vectorized operations in this modified function to speed up the process of warping each image to the output mosaic shape.

Next, I used the homography matrix and a translation matrix to warp each image to the correct position in the output mosaic. From there, I could warp the first image with just the translation transformation, and then warp the second image onto the output by combining the homography and translation transformations. To handle blending, I took a weighted average of the two images based on distance from the edge of each image, so that pixels closer to the edge would have a lower weight and pixels closer to the center would have a higher weight. This helped to smooth out edge artifacts, as seen in the mosaics below!

I had a lot of fun going around campus taking pictures of iconic landmarks at Berkeley (not pictured, Campanile shots) as well as exploring my photos from previous travels like my trip to Banff and Jasper National Park in Canada this summer!

Grimes 1




Source image 1 of Grimes Engineering Center

Grimes 2



Source image 2 of Grimes

Grimes Mosaic



Mosaic of Grimes

Inside Grimes 1




Source image 1 of inside Grimes

Inside Grimes 2



Source image 2 of inside Grimes

Inside Grimes Mosaic




Mosaic of Inside Grimes

Skywalk 1




Source image 1 of the Skywalk at Jasper National Park

Skywalk 2



Source image 2 of the Skywalk at Jasper National Park

Skywalk Mosaic




Mosaic of the Skywalk (a glass floor observation platform!)

Part B

B.1 Harris Corner Detection


To obtain the Harris corners for my images, I ran the sample code provided in our project spec and then ran `get_harris_corners()` on my input images. To visualize the corners overlaid on the image, I plotted the Harris corner coordinates as a scatterplot over the image. I decided to try this on my image sets of the Grimes Engineering Center, the Campanile, as well as a vertical shot I took of the World Trade Center (WTC) in NYC! When run on my images below, many corners were initially detected, with almost the entire space of the image being seemingly filled with corners.

Grimes Harris Corners




Harris corners detected on an image of Grimes

Campanile Harris Corners



Harris corners detected on an image of the Campanile


World Trade Center (WTC) Harris Corners



Harris corners detected on an image of the WTC!


However, Adaptive Non-Maximal Suppression (ANMS) was very helpful in narrowing down corners of interest from Harris corner detection. To implement ANMS, I followed the MOPS paper in initializing a suppression radius to compare corners to other corners. First, I retrieved the corner strength values for each corner, and then for each corner, I looped through all other corners to check if a corner was potentially stronger than the current corner, and then computed the radius between the two corners. If the radius was larger than the current suppression radius, then the larger suppression radius becomes the new radius. After completing the loop, we can take the top N corners (I believe the paper chose 500) with the largest suppression radius. As a result of keeping the top 500 corners with the largest suppression radii, we have a better distribution of corners across the image for better coverage in terms of image information.

Grimes ANMS Corners




ANMS corners detected on an image of Grimes

Campanile ANMS Corners



ANMS corners detected on an image of the Campanile

World Trade Center (WTC) ANMS Corners




ANMS corners detected on an image of the WTC!

B.2 Feature Descriptor Extraction


To implement feature descriptor extraction, I first sampled from a 40x40 window, but skipped over points that were too close to the edge of the image. From there, I resized down to 8x8 windows to extract patches in these dimensions. I then normalized the patches by subtracting the mean and dividing by the standard deviation. Finally, I flattened these patches into length = 64 vectors (to preserve all 8x8 values from our patch) to serve as our feature descriptors. Below I include 15 extracted feature descriptors for each of the below images. Essentially, each 8x8 patch for a corner can capture the structure of the image locally around the corner, which we can then leverage in matching features between corresponding images.

Grimes Feature Descriptors




8x8 Extracted Feature Descriptors from Grimes image

Campanile Feature Descriptors



8x8 Extracted Feature Descriptors from Campanile image

WTC Feature Descriptors



8x8 Extracted Feature Descriptors from WTC image


B.3 Feature Matching

Next, to match features from one image to the other, I essentially looped through each feature descriptor that we extracted from the previous section starting in image 1 and then calculated the difference in Euclidean distance between that descriptor and all descriptors in image 2. Then, we can find the smallest distance and then the next smallest distance, and apply Lowe's thresholding technique of dividing the smallest distance by the second smallest distance and comparing it to

a given threshold. The idea behind this is that a good match should have a small ratio (below our chosen threshold, I experimented with thresholds depending on the image, but more on that next!) because our theory is that the best match should be significantly better than the second best match, whereas bad matches may have multiple close contender distances.


To elaborate, for my images of Grimes and the Campanile, I found that a threshold of 0.7 worked well to filter out bad matches, whereas for my image of the WTC, I actually tried decreasing my threshold to 0.5 because I was worried some of my feature matches would skew my final mosaic, however, I found that this threshold was too strict and I wasn't getting enough matches. Below I include some diagrams of the features that were matched between using these different thresholds!

Grimes Feature Matching




Feature matches found for Grimes images

Campanile Feature Matching




Feature matches found for Campanile images

WTC Feature Matching (Threshold = 0.7)



Feature matches found for WTC images with higher threshold

WTC Feature Matching (Threshold = 0.5)



Feature matches found for WTC images with stricter threshold

B.4 RANSAC for Robust Homography


After determining our feature matches, we can improve the robustness of our homography estimation by implementing RANSAC. To do this, I looped through many iterations following the steps of:

1. Randomly sampling 4 pairs of points because 4 is the minimum number of points needed to compute a homography
2. Computing the homography matrix from these 4 pairs of points
3. Checking how well H explains the point pairs by calculating the difference in distance between the projected point and the actual point (I used L2 distance again here)
4. Keeping the best H (the one with the most inliers AKA best consensus)
5. Recomputing the H with all the inliers and excluding the outliers

Here, we can see how both the Grimes and Campanile mosaic combine images in the horizontal direction, whereas the WTC mosaic is the only vertical mosaic.


Overall, moving from part A to part B has been super interesting because we can experience the gratification of progressing from manually selecting points across our image sets to carefully identify correspondences to now being able to automatically detect corners, extract feature descriptors, match features, and then robustly compute homographies to automate the process of creating image mosaics!

Grimes Mosaic (with RANSAC)




Grimes mosaic using points from RANSAC

Campanile Mosaic (with RANSAC)



Campanile mosaic using points from RANSAC

WTC Mosaic (with RANSAC)



WTC mosaic using points from RANSAC